

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud!**DEPLOY NOW**

Kubernetes Objects Guide

August 11, 2022

KUBERNETES

[Home](#) » [DevOps and Development](#) » Kubernetes Objects Guide

Introduction

[Container orchestration](#) is an essential aspect of managing operational complexity in [DevOps](#). Teams adopting [Kubernetes](#) for containerized workload management aim to improve agility and speed by utilizing the platform's declarative approach and automation features.

One of the most important ways Kubernetes administrators interact with the platform is by creating and managing Kubernetes objects. Objects are instrumental in helping users to deploy apps and maintain the cluster.

This guide will provide a detailed overview of Kubernetes objects, analyze their structure, and offer helpful object management tips. It will also list and describe the most frequently used objects.

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud!phoenixNAP
GLOBAL IT SERVICES

DEPLOY NOW



Kubernetes Objects Guide



What are Kubernetes Objects?

Kubernetes objects are entities in a Kubernetes cluster that serve as a record of intent. Administrators create objects to express the cluster's desired state, and Kubernetes uses them to maintain this state automatically.

Another way to understand Kubernetes objects is by looking at them as class instances. Each created object references to a pre-defined class that tells the [API server](#) how to handle system resources and communicate with specific components.

To see the complete list of available objects in a Kubernetes cluster, type the following [kubectl command](#):

```
kubectl api-resources
```



marko@test-01: ~\$ kubectl api-resources

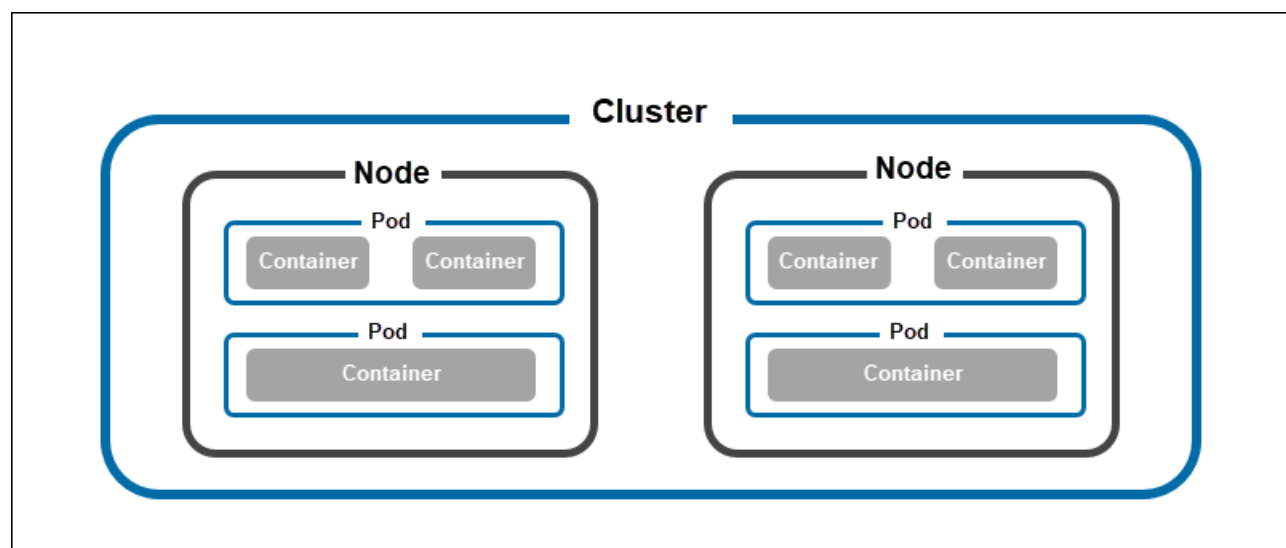
NAME	SHORTNAMES	APIVERSION	NAMESPACED	KIND
bindings		v1	true	Binding
componentstatuses	cs	v1	false	ComponentStatus
configmaps	cm	v1	true	ConfigMap
endpoints	ep	v1	true	Endpoints
events	ev	v1	true	Event
limitranges	limits	v1	true	LimitRange
namespaces	ns	v1	false	Namespace
nodes	no	v1	false	Node
persistentvolumeclaims	pvc	v1	true	PersistentVolumeClaim
persistentvolumes	pv	v1	false	PersistentVolume
Pods	po	v1	true	Pod
podtemplates		v1	true	PodTemplate
replicationcontrollers	rc	v1	true	ReplicationController
resourcequotas	quota	v1	true	ResourceQuota
secrets		v1	true	Secret
serviceaccounts	sa	v1	true	ServiceAccount
services	svc	v1	true	Service

DEPLOY NOW

The following sections explain the object concept by describing the most important Kubernetes objects.

Pods

[Kubernetes pods](#) are the smallest unit of deployment in Kubernetes. They reside on cluster nodes and have their IP addresses, enabling them to communicate with the rest of the cluster. A single pod can host one or more containers, providing storage and [networking resources](#).



One of the key characteristics of Kubernetes pods is that they are ephemeral. In practice, a pod can fail without impacting the system's functioning. Kubernetes automatically replaces each failed pod with a new pod replica and keeps the cluster running.

Aside from being container wrappers, pods also store configuration information that instructs Kubernetes on how to run the containers.

Services

Services provide a way to expose a **DEPLOY NOW** pods. Their purpose is to represent a set of pods that perform the same function and set the policy for accessing those pods.

Although pod failure is an expected event in a cluster, Kubernetes replaces the failed pod with a replica with a different IP address. This creates problems in communication between pods that depend on each other.

Using the **kube-proxy** process that runs on each cluster node, Kubernetes maps the service's virtual IP address to pod IP addresses. This process allows for easier internal networking but also enables exposing of the deployment to external networks via techniques such as [load balancing](#).

Volumes

Volumes are objects whose purpose is to provide storage to pods. There are two basic types of volumes in Kubernetes:

- **Ephemeral volumes** persist only during the lifetime of the pod they are tied to.
- [Persistent volumes](#), which are not destroyed when the pod crashes. Persistent volumes are created by issuing a request called **PersistentVolumeClaim** (PVC). Kubernetes uses PVCs to provision volumes, which then act as links between pods and physical storage.

Namespaces

The purpose of the Namespace object is to act as a separator of resources in the cluster. A single cluster can contain multiple namespaces, allowing administrators to organize the cluster better and simplify resource allocation.

A new cluster comes with multiple namespaces created for system purposes and the **default** namespace for users. Administrators can create any number of additional namespaces - for example, one for development and one for testing.

```
manko@test-main:~$ kubectl get namespaces
NAME                STATUS    AGE
default             Active   337d
development         Active   21s
kube-node-lease     Active   337d
kube-public          Active   337d
kube-system          Active   337d
testing             Active    8s
manko@test-main:~$
```

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud!

Note: Learn how to [Deploy a Kubernetes Namespace](#) or how to [Delete a Kubernetes Namespace](#) in our detailed guides.

[DEPLOY NOW](#)

Deployments

Deployments are controller objects that provide instructions on how Kubernetes should manage the pods hosting a [containerized application](#). Using deployments, administrators can:

- Scale the number of pod replicas.
- Rollout updated code.
- Perform rollbacks to older code versions.

Once created, the deployment controller monitors the health of the pods and nodes. In case of a failure, it destroys the failed pods and creates new ones. It can also bypass the malfunctioning nodes, enabling the application to remain functional even when a hardware error occurs.

ReplicationControllers

ReplicationControllers ensure that the correct number of pod replicas are running on the cluster at all times. When creating a ReplicationController, the administrator specifies the desired number of pods. The controller then maintains their number, creating additional pods and terminating the extra ones when necessary.

ReplicationControllers support equality-based selectors, which allow filtering by label keys and values. The controller will manage all the pods whose label matches the one provided in the **.spec.selector** field of the configuration file.

For example, the following declaration tells Kubernetes to run five **nginx** pods:

```
spec:
  replicas: 5
  selector:
    app: nginx
```



Since manually created pods are not automatically replaced when they fail, using replication is a recommended practice when the desired number of pods is one.

[DEPLOY NOW](#)

Note: Although they are not deprecated, ReplicationControllers are no longer the recommended way of setting up replication. Kubernetes documentation recommends using ReplicaSets in Deployments.

ReplicaSets

ReplicaSets serve the same purpose as ReplicationControllers, i.e. maintaining the same number of pod replicas on the cluster. However, the difference between these two objects is the type of selectors they support. While ReplicationControllers accept only equality-based selectors, ReplicaSets additionally support set-based selectors.

Set-based selectors allow using a set of values to filter keys. The statements accept three operators: **in**, **notin**, and **exists**. For example, the following **selector** section instructs the ReplicaSet to run an Nginx pod that belongs to **production** and **qa** environments:

```
selector:
  matchLabels:
    app: nginx
  matchExpressions:
    - {key: environment, operator: In, values: [production, qa]}
```

DaemonSets

DaemonSets are controller objects whose purpose is to ensure that specific pods run on specific (or all) nodes in the cluster. Kubernetes scheduler ignores the pods created by a DaemonSet, so those pods last for as long as the node exists. This object is particularly useful for setting up daemons that need to run on each node, like those used for cluster storage, log collection, and node monitoring.

By default, a `DaemonSet` creates a pod on every node in the cluster. If the object needs to target specific nodes, their selection is done using the `nodeSelector` field in the configuration file.

[DEPLOY NOW](#)

StatefulSets

While Deployments and Replication Controllers can handle stateless apps, stateful apps require a workload object called StatefulSet. A StatefulSet gives each pod a unique identity, which persists across pod restarts.

The applications that benefit from StatefulSet are those which require:

- **Unique persistent storage for each pod** - PersistentVolumeClaims that are used with Deployments provide shared storage for pod replicas. StatefulSets instead use VolumeClaimTemplates, which assign a unique PVC to each replica.
- **Unique network ID** - A headless service controls the pod network identity.
- **Ordered deployment, scaling, and rolling updates.**

Databases such as [MySQL](#) and [PostgreSQL](#) are examples of applications that are deployed using StatefulSets.

ConfigMaps

[ConfigMaps](#) are Kubernetes objects used to store container configuration data in key-value pairs. By separating configuration data from the rest of the container image, ConfigMaps enable the creation of lighter and more portable images. They also allow developers to use the same code with different configurations depending on whether the app is in the development, testing, or production phase.

A pod can be configured to use the ConfigMap data by mounting the ConfigMap as a volume inside the pod:

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud!

volumes:

- name: config

[DEPLOY NOW](#)

configMap

name: config-example

items:

- key: key1

path: config.example

Alternatively, the environment variables can be used pull the specific values from the configuration:

env:

- name: SPECIAL_LEVEL_KEY

valueFrom:

configMapKeyRef:

name: config-example

key: key1

Jobs

[Jobs](#) are workload controller objects that execute finite tasks. While other controller objects have the task of permanently maintaining the desired state and number of pods, jobs are designed to finish a task and terminate the associated pods. This property makes them useful for maintenance, monitoring, batch tasks, and work queue management.

Job instances run simultaneously or consecutively. Scheduled jobs are a separate controller object called [CronJob](#).



Note: Our Bare Metal Cloud solution offers [Rancher integration](#) for easy Kubernetes management. Get started by finding the [servers](#) that suit your needs.

Object Spec and Status

Object **spec** and object **status** are two nested fields in the object configuration that

Kubernetes utilizes to control the object.

DEPLOY NOW

- The **spec** field is used to declare the desired state of the object, i.e., the characteristics of the resource. The user provides this field when creating the object.
- The **status** field provides information about the current object state. This field is provided by Kubernetes during the lifetime of the object.

Kubernetes control plane monitors the status of every object in the cluster and attempts to match the current state to the desired state.

For example, consider the following deployment YAML:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-test
  labels:
    app: nginx
spec:
  replicas: 5
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.23.1
        ports:
        - containerPort: 80
```

The **spec** field states that the desired number of replicas for the nginx deployment is five. After the deployment is successfully created, Kubernetes monitors its status and updates

the status field accordingly. If one replica fails, the status field reports only four running replicas, which triggers Kubernetes

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud!

DEPLOY NOW



Object Required Fields

When the user wants to create a Kubernetes object, the following fields must be provided in the YAML file:

- **apiVersion** - Specifies the version of Kubernetes API for creating the object.
- **kind** - Provides the object type, for example, Deployment, ReplicaSet, or Service.
- **metadata** - Lists object identifiers, such as its name, UID, labels, and namespace.
- **spec** - States the desired state for the object, like the number of replicas and the container image.

Object Names and IDs

Each Kubernetes object has a name and a UID. They are used to identify the object across the cluster.

Object names are user-defined and unique for object types. This means that, for example, there can be only one pod named **test-app** within a single namespace, but the deployment of the same name is allowed to exist.

The system assigns object UIDs, which are unique for each object instance across types and namespaces. This means that even the historical versions of the same object have different UIDs.

Working with Kubernetes Objects

Kubernetes objects are managed using various GUI dashboards or using the kubectl CLI tool. With kubectl, users can manage objects by employing three distinct management techniques:

- **Imperative commands** are single action words that can change the cluster in one step. However, tracking changes in the cluster is significantly more difficult if users apply many of these commands.
- **Imperative configuration** specifies a command (such as **create**), options, and a YAML or JSON file containing object configuration. This approach allows storing the configuration in source control systems (such as [Git](#)) for easier access but features a

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud!

- **Declarative configuration** transfers the necessary operations to kubectl while requiring the user only to provide configuration files. When the user applies a configuration file, kubectl performs the necessary actions on the object.

DEPLOY NOW



Note: Always use a single technique to manage a single Kubernetes object. Mixing techniques can have undesired consequences.

Create Objects

You can create Kubernetes using any of the management techniques listed above.

For example, the following imperative command creates an Nginx deployment:

```
kubectl create deployment nginx --image nginx
```

The same deployment can be created using the imperative object configuration:

```
kubectl create -f nginx.yaml
```

The third option involves using the declarative approach:

```
kubectl apply -f nginx.yaml
```

Edit Objects

Kubectl features multiple subcommands for editing an existing object. Below are the most common commands:

- **scale** - Update the replica count of the controller to perform horizontal scaling.
- **annotate** - Edit the annotation of the object.
- **label** - Edit the object label.

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud! 

The commands above allow users [DEPLOY NOW](#) without having to know the specific fields that need to change. The following set of commands requires a better understanding of the object schema:

- **edit** - Open and edit the object's configuration in a text editor.
- **patch** - Use a patch string to modify specific object fields.

Delete Objects

Use the **kubectl delete** command to delete objects from the cluster. To employ the imperative approach with this command, pass the object as a command argument. For example:

```
kubectl delete deployment/nginx
```

You can also delete an object by using imperative object configuration and providing the configuration file.

```
kubectl delete -f nginx.yaml
```

and make technology less daunting to everyone.

Conclusion

Next you should read

This guide introduced you to the concept of a Kubernetes object. It presented the most commonly used objects and provided insight into their form and function. Finally, some

DevOps and Kubernetes Architecture with Diagrams if you would like to delve

Development,
Virtualization

Understanding

Kubernetes

Architecture with

Diagrams

November 12, 2019

ues were listed and explained.

Yes

No

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud!



DEPLOY NOW

[DevOps and
Development](#)
**Kubernetes Use
Cases**

August 11, 2022

This article focuses on Kubernetes as the most popular orchestration tool on the market. We will analyze a broad range of Kubernetes use cases from the platform's...

READ MORE

[DevOps and
Development,
Networking](#)
**Kubernetes
Networking Guide**

June 7, 2022

The guide shows the essentials of Kubernetes networking

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud!**DEPLOY NOW**

DevOps and
Development,
Networking

Kubernetes

**Service Discovery
Guide**

June 2, 2022

A service helps manage internal and external traffic to pods through IP addresses, ports, and DNS records. Service discovery...

READ MORE

ar

io

u

s

Live Chat

Get a Quote

Support | 1-855-330-1509

Sales | 1-877-588-5918

u

b

er

n

et

e

s.

..

R

Get 15 TB FREE bandwidth (5 TB in Singapore) with Bare Metal Cloud!



A

Privacy Center

Do

DEPLOY NOW

Personal information

Contact Us

Legal **M**

Privacy Policy

Terms of Use

R
DMCA

E
GDPR

Sitemap

© 2022 Copyright phoenixNAP | Global IT Services. All Rights Reserved.